

Java Programming Course

Parallel Programming with Java



Faculty of Information Technologies
Industrial University of Ho Chi Minh City

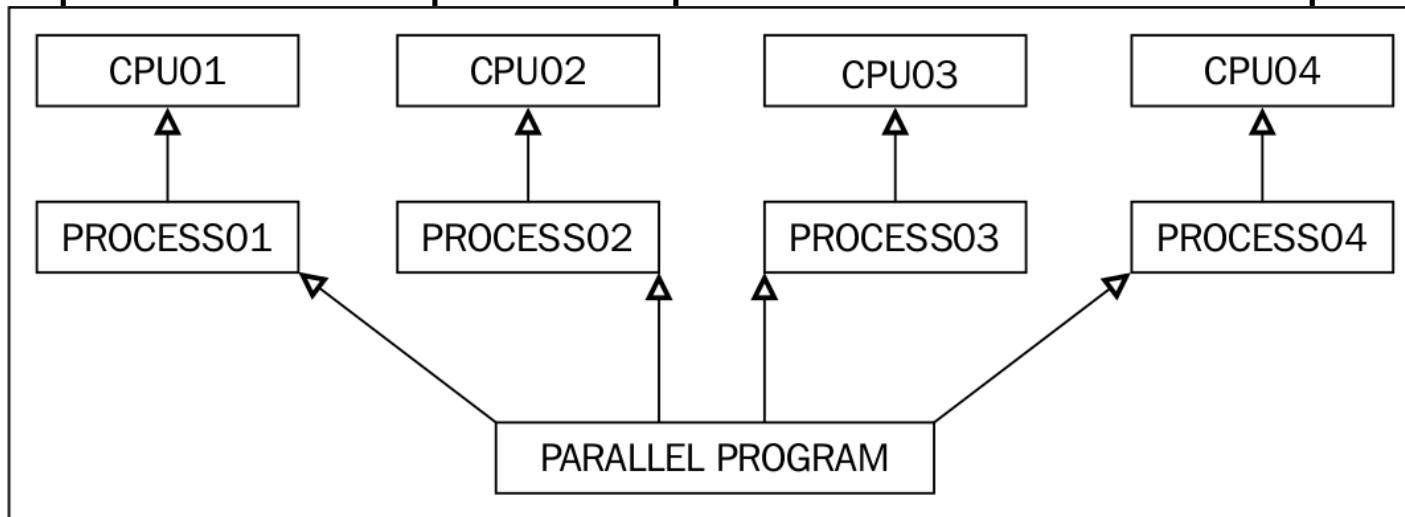
Session objectives

- Introduction
- Fork/Join Framework
 - ForkJoinPool
 - ForkJoinTask
 - RecursiveTask
 - RecursiveAction
- Task Parallelism
- Functional Parallelism
- Loop Parallelism
- Pipelining Parallelism



Introduction

- ◇ Parallel Programming is a process of breaking down a complex problem into smaller and simpler tasks, which can be executed at the same time using a number of computer resources.
- ◇ In the process of Parallel Programming, the tasks which are independent of each other are executed parallelly using different computers or multiple cores present in a CPU of a computer.





Introduction (cont'd)

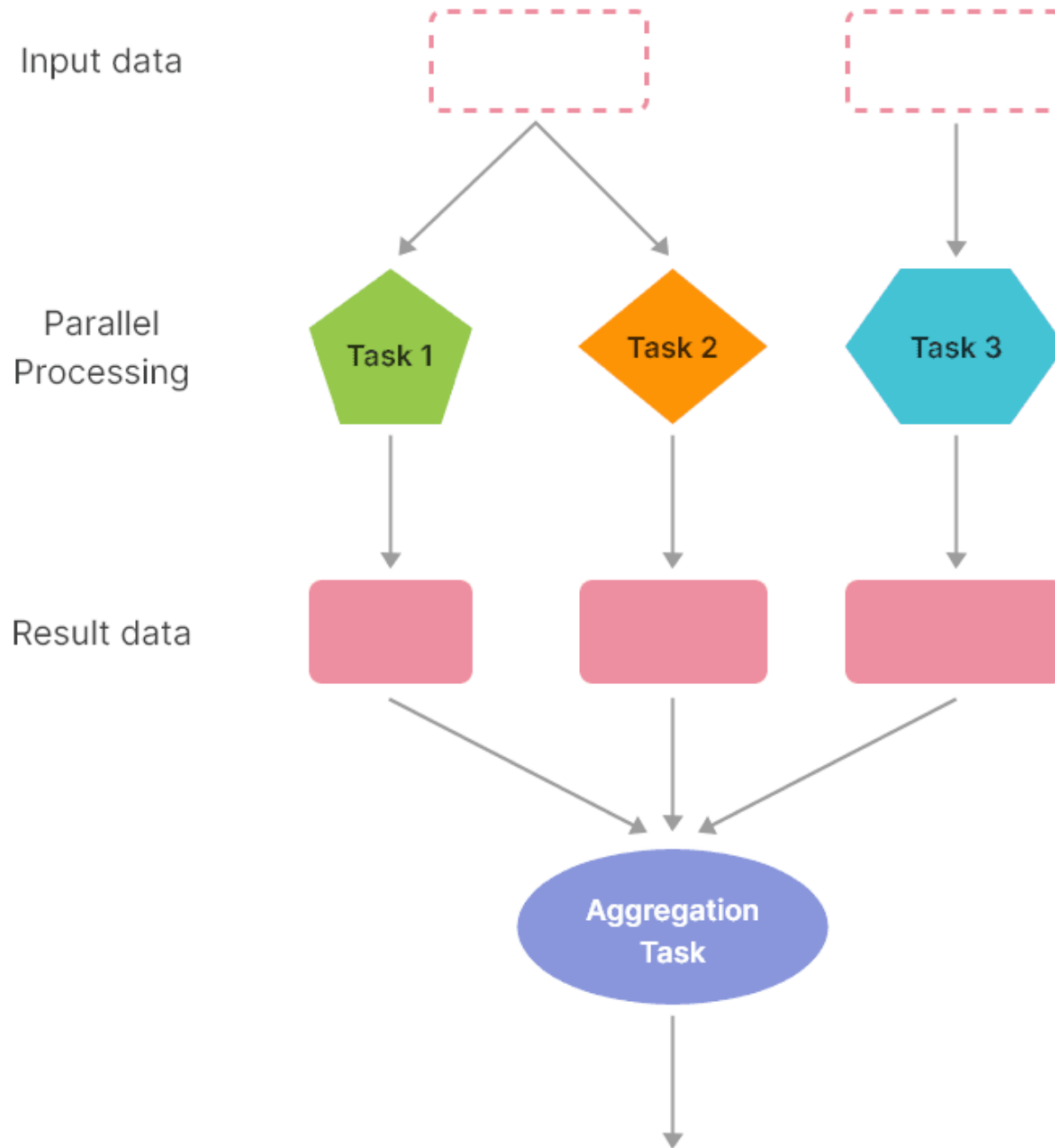
- ◇ Parallelism is a form of computation in which many instructions are carried out simultaneously, operating on the principle that large problems can often be divided into smaller ones, which are then solved concurrently (in parallel).
- ◇ Styles of parallelism:
 - Task parallelism
 - Data parallelism



Task Parallelism

- ◇ Also known as function parallelism.
- ◇ Approach that focuses on distributing tasks rather than data across different processing units. Each of these tasks can be a separate function or a method operating on different data or performing different computations. This type of parallelism is a great fit for problems where different operations can be performed concurrently on the same or different data.

Task Parallelism (cont'd)

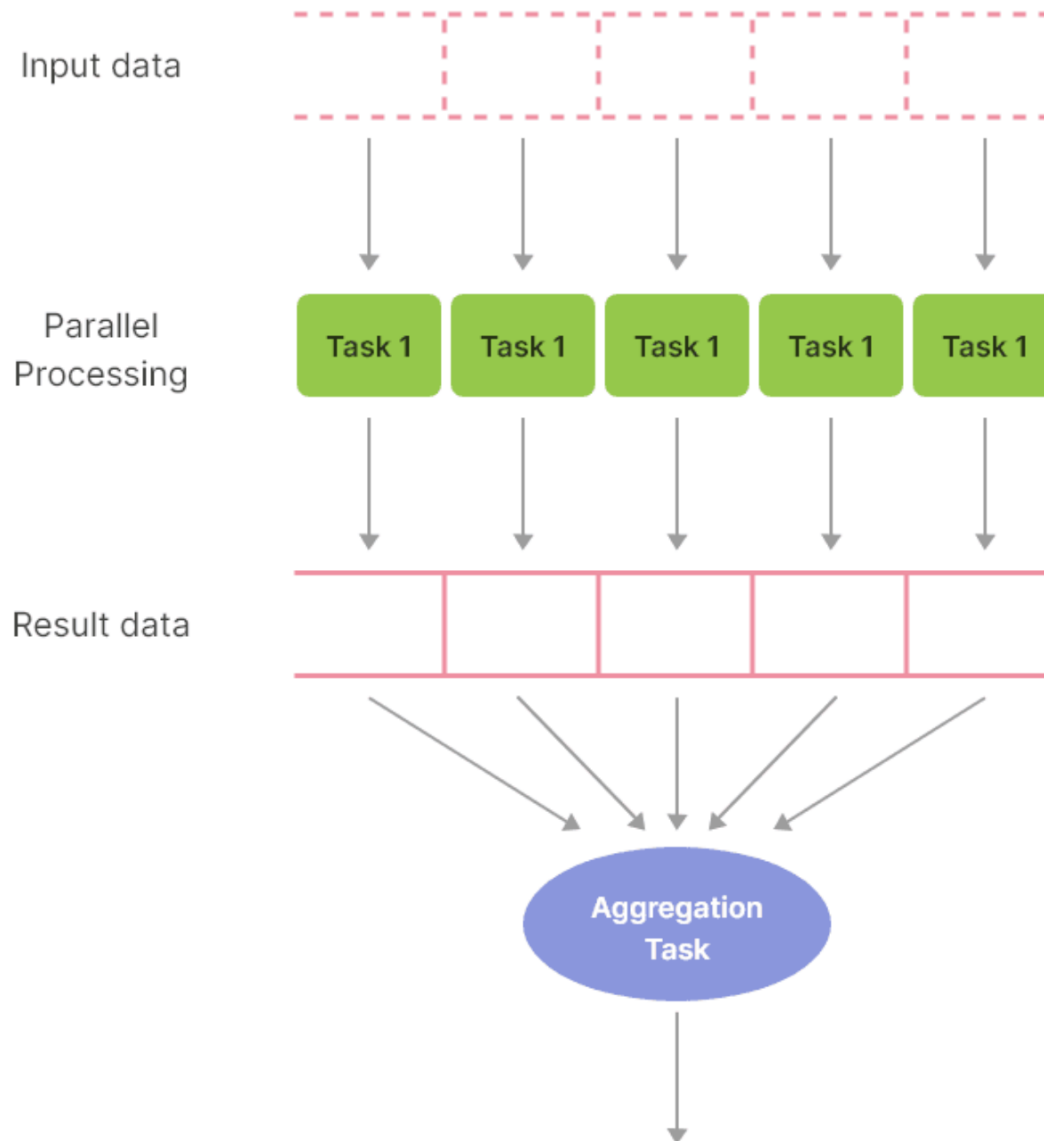




Data Parallelism

- ◇ The dataset is split into smaller, more manageable chunks, or partitions. Each partition is processed independently by separate tasks running the same operation. This distribution is done in a way that each task operates on a different core or processor, enabling high-level parallel computation.

Data Parallelism (cont'd)



Why is Parallel Programming used?

- ◇ Basically, Parallel Programming refers to the parallel execution of processes due to the availability of multiple resources such as processing cores.
- ◇ Parallel Programming is considered to be a much more efficient method than multithreading. The parallel tasks which are solved simultaneously are combined to give the final solution to a larger problem.
- ◇ The Java Standard Environment (Java SE) provides a programmer with the " Fork/Join Framework ". The Fork/Join Framework helps a programmer to implement Parallel Programming easily in their applications.

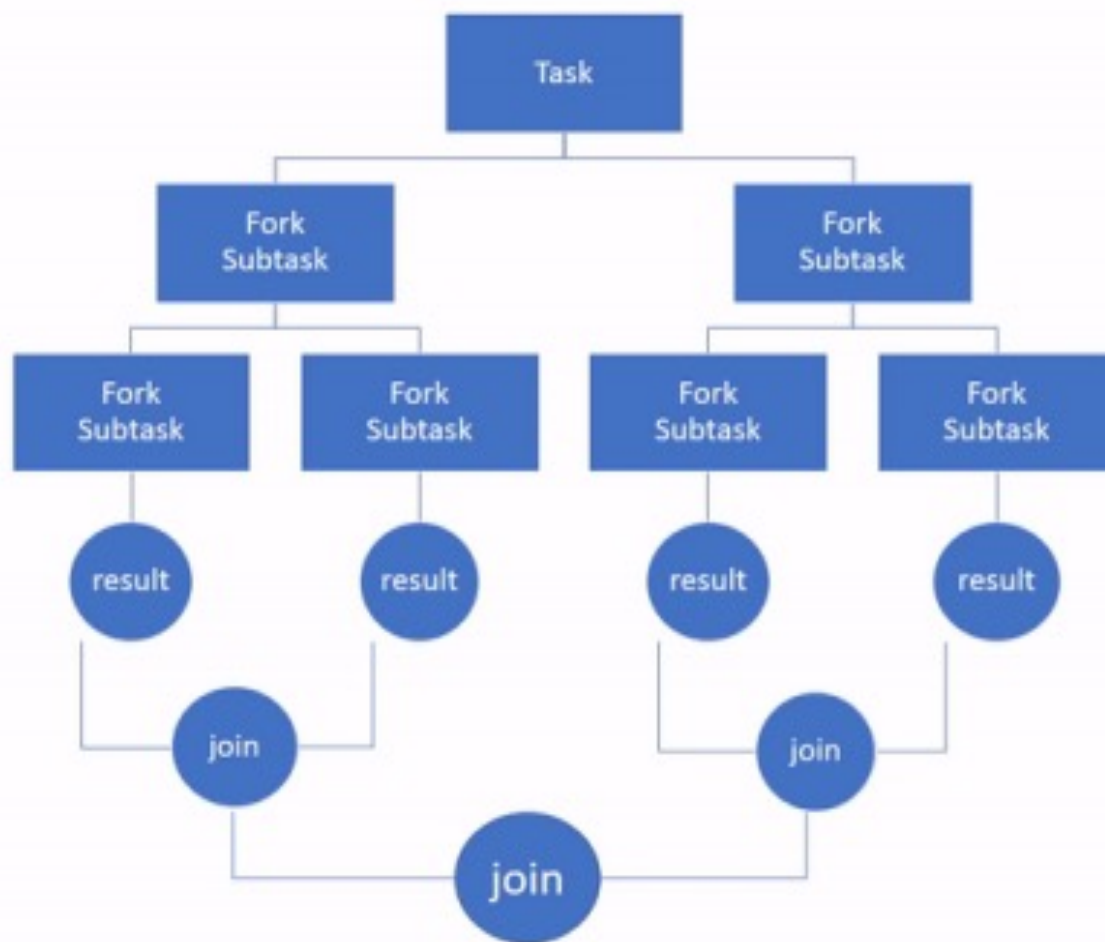
Fork/Join Framework

An overview of Java ForkJoinPool

- ◇ Since Java version 1.7
- ◇ The ForkJoinPool is similar to the Java ExecutorService but with one difference that it easily split the working task into smaller tasks which are then submitted to the ForkJoinPool (this pool manages worker threads of type ForkJoinWorkerThread); and finally joins the results (if any).

Fork/Join Framework

An overview of Java ForkJoinPool (1)



Fork/Join Framework

An overview of Java ForkJoinPool (2)

- ◇ The fork-join pool supports a style of parallel programming that solves problems by “divide & conquer”.
- ◇ Worker threads can execute only one task at a time, but the ForkJoinPool doesn't create a separate thread for every single subtask. Instead, each thread in the pool has its own double-ended queue (or deque) that stores tasks.
- ◇ This architecture is vital for balancing the thread's workload with the help of the work-stealing algorithm.

Fork/Join Framework

An overview of Java ForkJoinPool (3)

solve(problem):

if problem is small enough:

 solve problem directly (sequential algorithm)

else:

 for part in subdivide(problem)

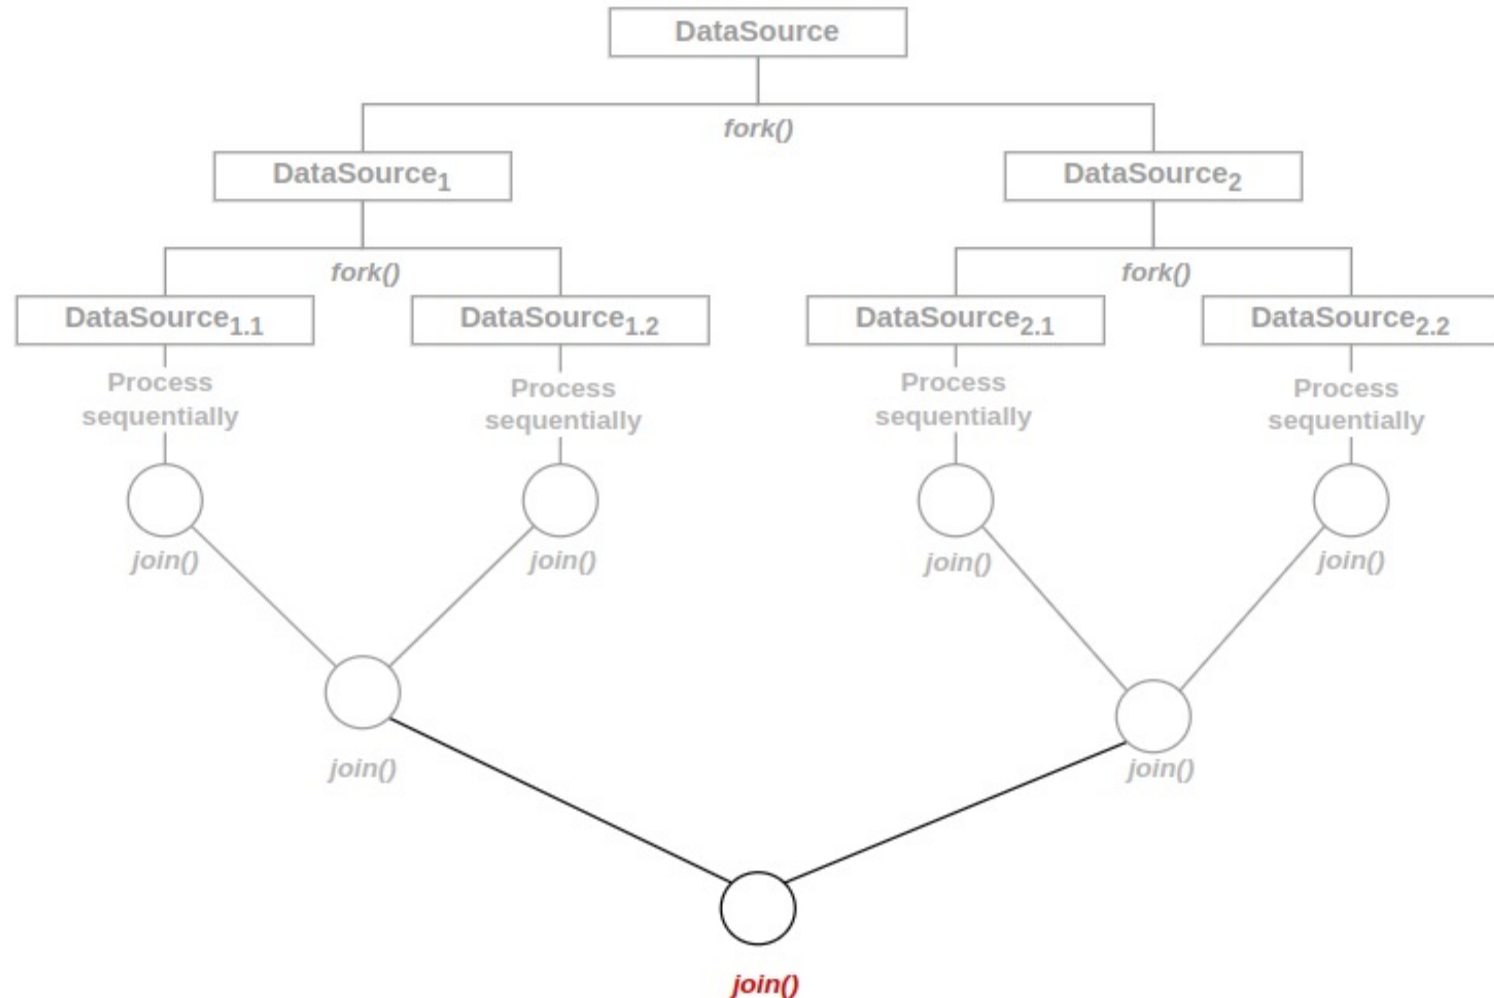
 fork subtask to **solve(part)**

 join all subtasks spawned in previous loop

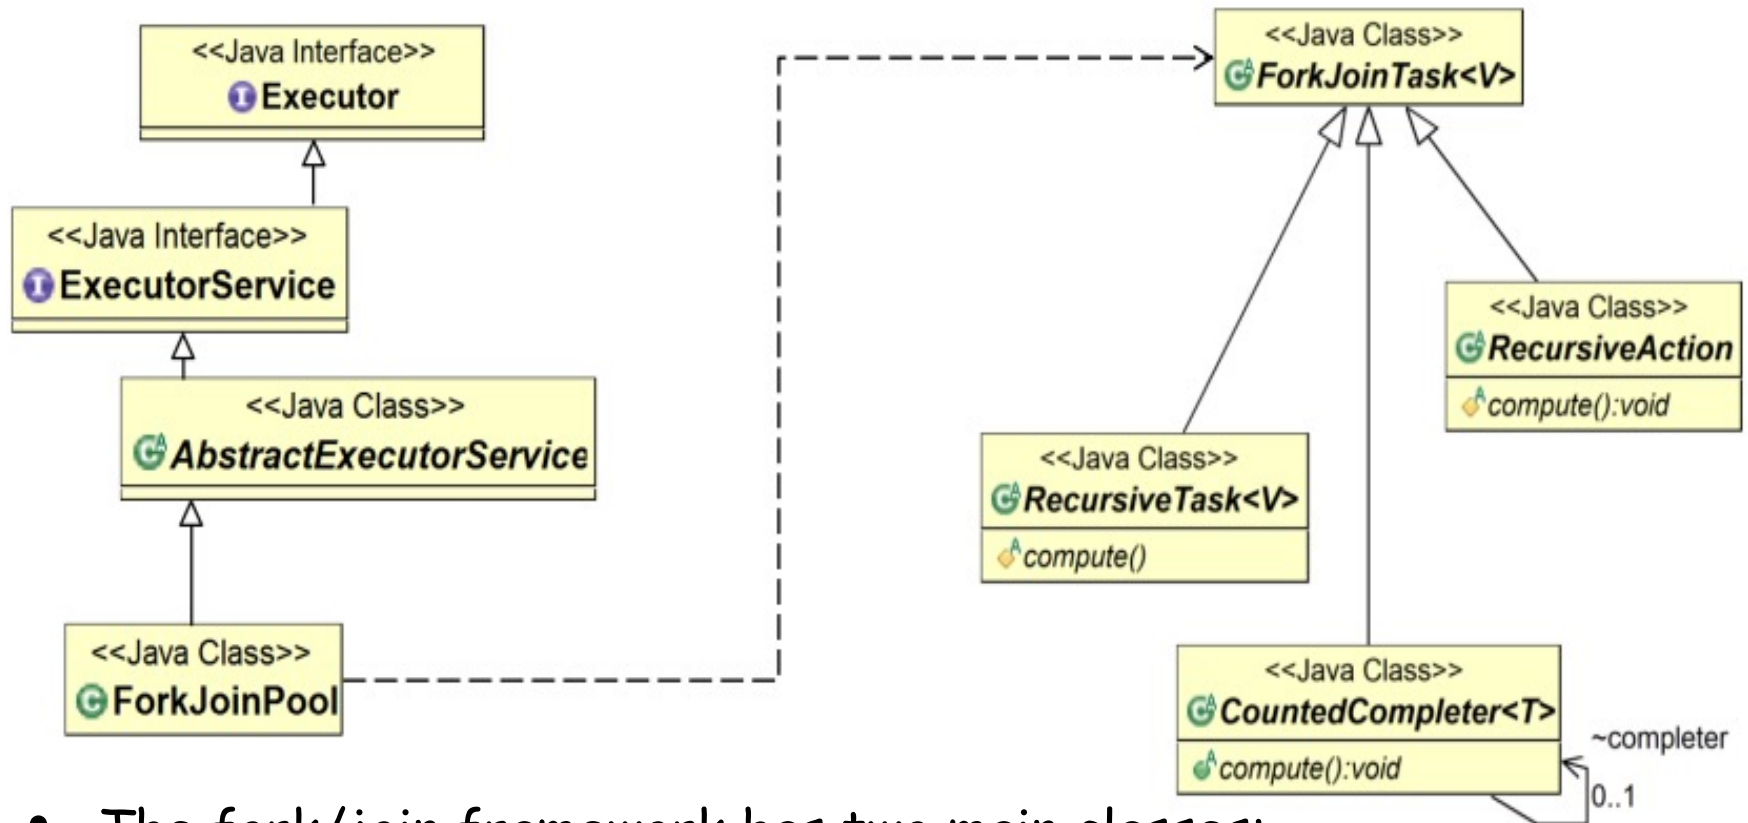
 return combined results

Fork/Join Framework

An overview of Java ForkJoinPool (4)



Fork/Join Framework Structure & Functionality



- The fork/join framework has two main classes:
 - ForkJoinPool
 - ForkJoinTask

Fork/Join Framework

ForkJoinPool class

- ◇ Is an implementation of the interface `ExecutorService`.
- ◇ In general, executors provide an easier way to manage concurrent tasks than plain old threads. The main feature of this implementation is the work-stealing algorithm.
- ◇ There's a common `ForkJoinPool` instance available to all applications that you can get with the static method `commonPool()`
 - `ForkJoinPool commonPool = ForkJoinPool.commonPool();`
- ◇ Constructors:
 - `ForkJoinPool()`
 - `ForkJoinPool(int parallelism)`
 - ...

Fork/Join Framework

ForkJoinPool class (1)

◇ Key methods:

- `execute()`: Desired asynchronous execution; call its fork method to split the work between multiple threads.
- `invoke()`: Await to obtain the result; call the invoke method on the pool.
- `submit()`: Returns a Future object that you can use for checking the status and getting the result on its completion.

Fork/Join Framework

ForkJoinTask class

- ◇ ForkJoinPool class invokes a task of type ForkJoinTask, which you have to implement by extending one of its two subclasses:
 - **RecursiveAction**: which represents tasks that do not yield a return value, like a Runnable.
 - **RecursiveTask**: which represents tasks that yield return values, like a Callable.
- ◇ ForkJoinTask subclasses also contain the following methods:
 - **fork()**: which allows a ForkJoinTask to be scheduled for asynchronous execution (launching a new subtask from an existing one).
 - **join()**: which returns the result of the computation when it is done, allowing a task to wait for the completion of another one.

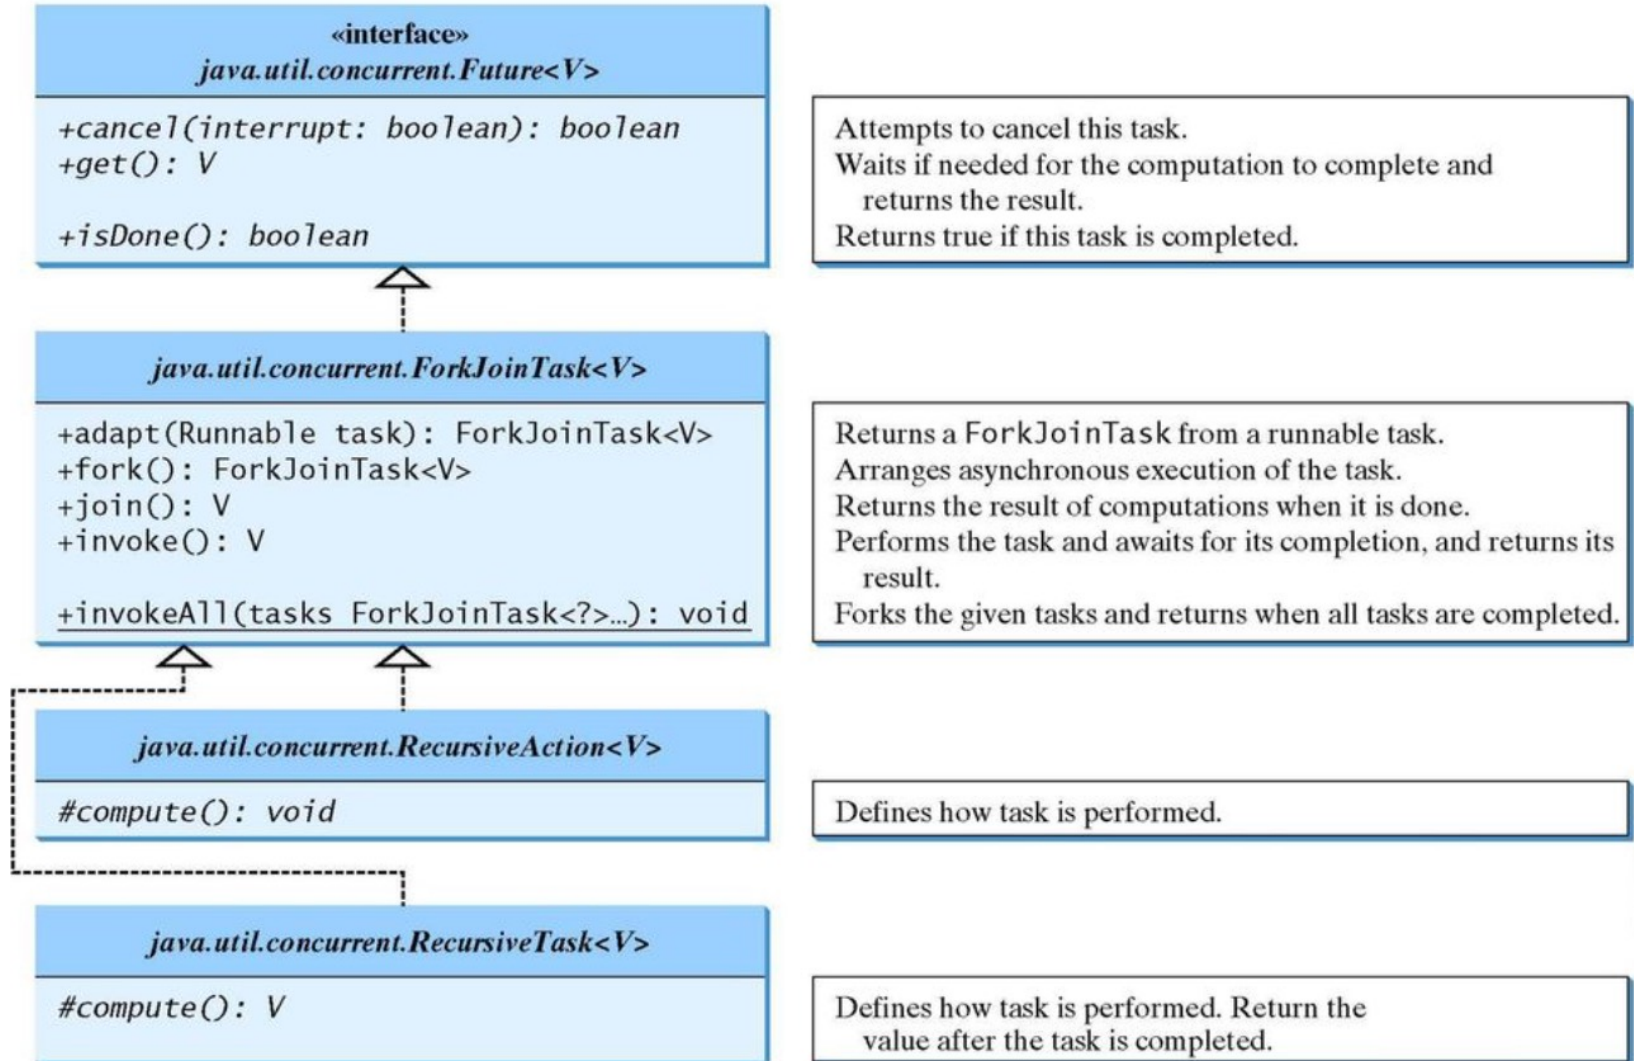
Fork/Join Framework

ForkJoinTask class (1)

- ◇ The ForkJoinTask class provides several methods for checking the execution status of a task.
 - `isDone()`: returns true if a task completes in any way.
 - `isCompletedNormally()`: returns true if a task completes without cancellation or encountering an exception.
 - `isCancelled()`: returns true if the task was cancelled.
 - `isCompletedabnormally()`: returns true if the task was either cancelled or throw an exception.

Fork/Join Framework

ForkJoinTask class (1)



Fork/Join Framework

Standard setup code

1. Create a ForkJoinPool (create first thread)
 - `ForkJoinPool pool = ForkJoinPool.commonPool();`
2. Do subclass `RecursiveTask<V>` or `RecursiveAction`
 - `SumArray` extends `RecursiveAction`
 - `SumArray` extends `RecursiveTask<V>`
3. Do override `compute()`
4. Do return a `V` from `compute`
5. Do call `fork()` -> Start a new thread
6. Do call `join (which returns answer)` -> Wait for a thread to finish
7. Do call `compute()` to hand-optimize
8. Do create a pool and call `invoke`
 - `SumArray task = new SumArray(data);`
 - `pool.invoke(task);`

Task Parallelism

Creating Tasks with Fork/Join FW

- ◇ Tasks are the most basic unit of parallel programming,
- ◇ Modern task-based approaches for parallel programming:
 - Task creation
 - Task terminations
 - And the “computation graph”: work, span, ideal parallelism, ...
- ◇ In this framework, a task can be specified in the `compute()` method of a user-defined class that extends the standard `RecursiveAction` class in the FJ framework.

Task Parallelism

Example - SumAction

```
private static class SumAction extends RecursiveAction {
    private static final int SEQUENTIAL_THRESHOLD = 5_000_000;
    private List<Long> data;
    private long total = 0L;

    private long getTotal() { return total; }

    public SumAction(List<Long> data) { this.data = data; }

    @Override
    protected void compute() {
        if(data.size() <= SEQUENTIAL_THRESHOLD) {
            total = computeSummary();
        } else {
            // Recursive case: Calculate new range
            int mid = data.size() / 2;
            SumAction firstSubTask = new SumAction(data.subList(0, mid));
            SumAction secondSubTask = new SumAction(data.subList(mid, data.size()));

            // Queue the first task
            firstSubTask.fork();

            // Return the sum of all subtasks
            // compute the second task
            // wait for the first task result
            secondSubTask.compute();
            firstSubTask.join();

            // or simply call
            // invokeAll(firstSubtask, secondSubtask);

            total = firstSubTask.getTotal() + secondSubTask.getTotal();
        }
    }

    /** Method that calculates the sum */
    private long computeSummary() {
        long sum = 0L;
        for (Long l: data) {
            sum += l;
        }

        return sum;
    }
}
```


Task Parallelism

Computation Graphs

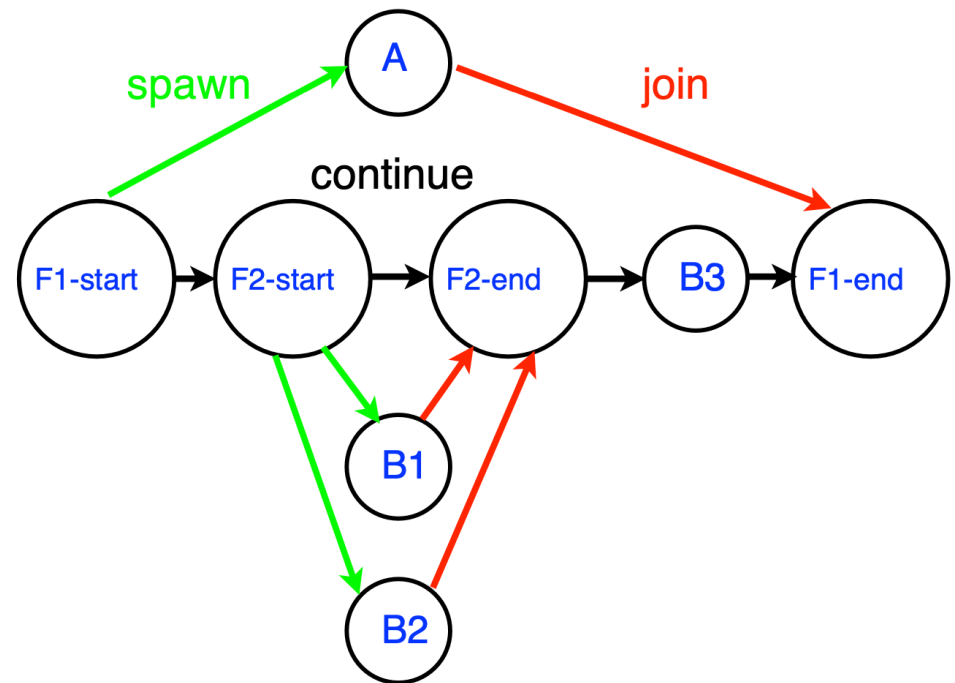
- ◇ A Computation Graph (CG) captures the dynamic execution of a parallel program, for a specific input
- ◇ CG nodes are "steps" in the program's execution
 - A step is a sequential sub-computation without any spawned, begin-finish or end-finish operations
- ◇ CG edges represent ordering constraints
 - "Continue" edges define sequencing of steps within a task
 - "Spawn/fork" edges connect parent tasks to child spawned tasks
 - "Sync/Join" edges connect the end of each spawned task to all join operations on that task.
- ◇ All computation graphs must be acyclic
 - It is not possible for a node to depend on itself
- ◇ Computation graphs are examples of "directed acyclic graphs" (DAGs)

Task Parallelism

Computation Graphs (1)

1. `must finish { //F1`
2. `spawn { A; }`
3. `must finish { //F2`
4. `spawn { B1; }`
5. `spawn { B2; }`
6. `} //F2`
7. `B3;`
8. `} //F1`

Computation Graph



Key idea: if two statements, X and Y have **no path of directed edges** from one to the other, then they can run in parallel with each other.



Functional Parallelism

- ◇ Will learn about futures, stream, ...
- ◇ Learn Java API for functional parallelism including Fork/Join framework and the Stream API's



Functional Parallelism

Creating Tasks with Fork/Join FW

- ◇ Extends the RecursiveTask
- ◇ `compute()` method return type (non-void)
- ◇ A method call like `left.join()` waits for the task referred to by object `left` in both cases, but also provides the task's return value in the case of future tasks.

Functional Parallelism

Example SumTask

```
class SumTask extends RecursiveTask<Long> {  
    private static final int SEQUENTIAL_THRESHOLD = 5_000_000;  
  
    private List<Long> data;  
  
    public SumTask(List<Long> data) {  
        this.data = data;  
    }  
  
    @Override  
    protected Long compute() {  
        if(data.size() <= SEQUENTIAL_THRESHOLD) {  
            long sum = computeSummary();  
            System.out.format("Sum of %s: %d\n", data.toString(), sum);  
            return sum;  
        } else {  
            // Recursive case  
            // Calculate new range  
            int mid = data.size() / 2;  
            SumTask firstSubTask = new SumTask(data.subList(0, mid));  
            SumTask secondSubTask = new SumTask(data.subList(mid, data.size()));  
  
            // Queue the first task  
            firstSubTask.fork();  
  
            // Return the sum of all subtasks  
            // compute the second task  
            // wait for the first task result  
            return secondSubTask.compute() + firstSubTask.join();  
        }  
    }  
}
```

Functional Parallelism

Java Stream

- ◇ Java Stream can provide a functional approach to operating on collections of data.
- ◇ As an example, the following pipeline can be used to compute the average age of all active students using Java streams:

```
students.stream()  
    .filter(s -> s.checkIsCuurent())  
    .mapToDoblue(a -> a.getAge())  
    .average();
```

Functional Parallelism

Java Stream

- ◇ An important benefit of using Java streams when possible is that the pipeline can be made to execute in parallel by designating the source to be a parallel stream:
 - simply replacing:
`students.stream()`
 - in the above code by:
`students.parallelStream()`
or
`Stream.of(students).parallel()`

Loop Parallelism

Parallel Loops

- ◇ The most general way is to think of each iteration of a parallel loop as an async task, with a finish construct encompassing all iterations

```
finish {  
    for (p = head; p != null ; p = p.next)  
        async compute(p);  
}
```

```
-> forall (i : [0:n-1]) a[i] = b[i] + c[i]
```

- ◇ Java streams can be an elegant way of specifying parallel loop computations

```
a = IntStream.rangeClosed(0, N-1).parallel().toArray(i  
-> b[i] + c[i]);
```

FAQ





That's all for this session!

Thank you all for your attention and patient !